

## Question 1

- lm:

```
lm(Y ~ ., data = train_df)
```

- Stepwise:

```
initial <- lm(formula = Y ~ 1, data = train_df)
final <- lm(formula = Y ~ ., data = train_df)
step(object=initial, scope=list(upper=final))
```

- Ridge:

```
lm.ridge(Y ~ ., lambda = seq(0,100,.05), data=train_df)
```

- Lasso:

```
cv.glmnet(x=train_matrix[, -1], y = train_df[, 'Y'])
```

- PLS:

```
plsr(Y ~ ., data = train_df, validation = "CV")
```

- Random Forest:

```
randomForest(Y ~ ., data=train_df, mtry = 11,
              nodesize = 15, ntree = 2000)
```

- Boosting:

```
gbm(Y ~ ., data = train_df, distribution = "gaussian",
     n.trees = 10000, interaction.depth = 8,
     shrinkage = 0.01)
```

- Neural Network:

```
nnet(y = Y.train, x = X.train, linout = TRUE, size = 31,
     decay = 0.8, maxit = 500)
```

## Question 2

We used 5 replicates of 10-folds CV on each model and computed the corresponding MSPEs. By doing a boxplot on the relative MSPE, we chose our final Boosting model because it outperformed other models almost all of the time.

## Question 3

(a) Random forest, Boosting and Neural Net all have tuning parameters. We used 5 replicates of 10-folds CV to tune each of them on a grid of the parameters. Then compare the models based on their computed RMSPEs.

(b)

- Random forest: we tuned on the combination of  $n_{\text{tree}} = (500, 1000, 1500, 2000)$ ,  $m_{\text{try}} = 3, 4, 5, \dots, 12$ , and  $\text{nodesize} = 2, 3, 4, \dots, 21$
- Boosting: we tuned on the combination of  $\text{shrinkage} = (0.0001, 0.001, 0.01, 0.1)$ ,  $\text{interaction.depth} = (3, 4, 5, 6, 7, 8, 9, 10)$ , and  $n_{\text{trees}} = (2000, 5000, 10000, 150000)$
- Neural Net: we tuned on the combinations of  $\text{size} = 1, 2, 3, \dots, 200$  and  $\text{decay} = 0, 0.1, 0.2, \dots, 1$

## Question 4

We choose Boosting with tuned parameter as our prediction machine.

```
library(gbm)
# Read in test and train data
rm(list = ls())
train_df = read.csv("Data/Data2020.csv")
test_df = read.csv("Data/Data2020testX.csv")

set.seed(1)
gbm.model = gbm(Y ~ ., data = train_df, distribution = "gaussian",
                 n.trees = 10000, interaction.depth = 8, shrinkage = 0.01)
n.trees.best = gbm.perf(gbm.model, plot.it = F) * 2 # Number of trees
prediction = predict(gbm.model, test_df, n.trees.best)
write.table(prediction, "Final/prediction.csv", sep = ",",
            row.names = F, col.names = F)
```

## 1 Question 5(Bonus)

We think X2, X4, X8, X12 are important